

email-spam-classification

June 20, 2024

```
[168]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import nltk
```

```
[169]: data=pd.read_csv('spam.csv', encoding='ISO-8859-1')
data.head()
```

```
[169]:      v1      v2 Unnamed: 2 \
0  ham  Go until jurong point, crazy.. Available only ...      NaN
1  ham              Ok lar... Joking wif u oni...      NaN
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...      NaN
3  ham  U dun say so early hor... U c already then say...      NaN
4  ham  Nah I don't think he goes to usf, he lives aro...      NaN

      Unnamed: 3 Unnamed: 4
0      NaN      NaN
1      NaN      NaN
2      NaN      NaN
3      NaN      NaN
4      NaN      NaN
```

Data Cleaning

```
[170]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   v1          5572 non-null   object
1   v2          5572 non-null   object
2   Unnamed: 2  50 non-null     object
3   Unnamed: 3  12 non-null     object
4   Unnamed: 4  6 non-null      object
dtypes: object(5)
```

memory usage: 217.8+ KB

```
[171]: # dropping unnecessary columns
data.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1, inplace=True)
```

```
[172]: data.head()
```

```
[172]:      v1                                     v2
0   ham  Go until jurong point, crazy.. Available only ...
1   ham                                Ok lar... Joking wif u oni...
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...
3   ham  U dun say so early hor... U c already then say...
4   ham  Nah I don't think he goes to usf, he lives aro...
```

```
[173]: # renaming columns
data=data.rename(columns={'v1': 'target', 'v2': 'Email'})
data.head()
```

```
[173]:      target                                     Email
0   ham  Go until jurong point, crazy.. Available only ...
1   ham                                Ok lar... Joking wif u oni...
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...
3   ham  U dun say so early hor... U c already then say...
4   ham  Nah I don't think he goes to usf, he lives aro...
```

```
[174]: data.isna().sum()
```

```
[174]: target      0
Email        0
dtype: int64
```

```
[175]: data.duplicated().sum()
```


```
[175]: 403
```

```
[176]: # Dropping duplicate values
data=data.drop_duplicates()
```

EDA

```
[177]: labels=data['target'].unique()
```

```
[178]: plt.figure(figsize=(4,3))
plt.pie(x=data['target'].value_counts(), labels=labels, autopct='%1.1f%%:')
plt.show()
```



```
[179]: # checking total characters in each email
data['num_characters']=data['Email'].apply(len)
```

```
[180]: # checking for no of words in each email
data['num_words']=data['Email'].apply(lambda x:len(nltk.word_tokenize(x)))
```

```
[181]: # checking for no of sentences in each mail
data['num_sentences']=data['Email'].apply(lambda x:len(nltk.sent_tokenize(x)))
```

```
[182]: data.head()
```

```
[182]:   target      Email  num_characters  \
0    ham  Go until jurong point, crazy.. Available only ...      111
1    ham                Ok lar... Joking wif u oni...       29
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...     155
3    ham  U dun say so early hor... U c already then say...     49
4    ham  Nah I don't think he goes to usf, he lives aro...     61
```

```
   num_words  num_sentences
0         23             2
1          8             2
2         37             2
3         13             1
4         15             1
```

```
[183]: # stats for not spam Emails
data[data['target']=='ham'].describe()
```

```
[183]:   num_characters  num_words  num_sentences
count    4516.000000  4516.000000    4516.000000
mean         70.459256    16.957484     1.815545
```

std	56.358207	13.394052	1.364098
min	2.000000	1.000000	1.000000
25%	34.000000	8.000000	1.000000
50%	52.000000	13.000000	1.000000
75%	90.000000	22.000000	2.000000
max	910.000000	219.000000	38.000000

```
[184]: # stats for spam Emails
data[data['target']=='spam'].describe()
```

```
[184]:
```

	num_characters	num_words	num_sentences
count	653.000000	653.000000	653.000000
mean	137.891271	27.474732	2.969372
std	30.137753	6.893007	1.488910
min	13.000000	2.000000	1.000000
25%	132.000000	25.000000	2.000000
50%	149.000000	29.000000	3.000000
75%	157.000000	32.000000	4.000000
max	224.000000	44.000000	9.000000

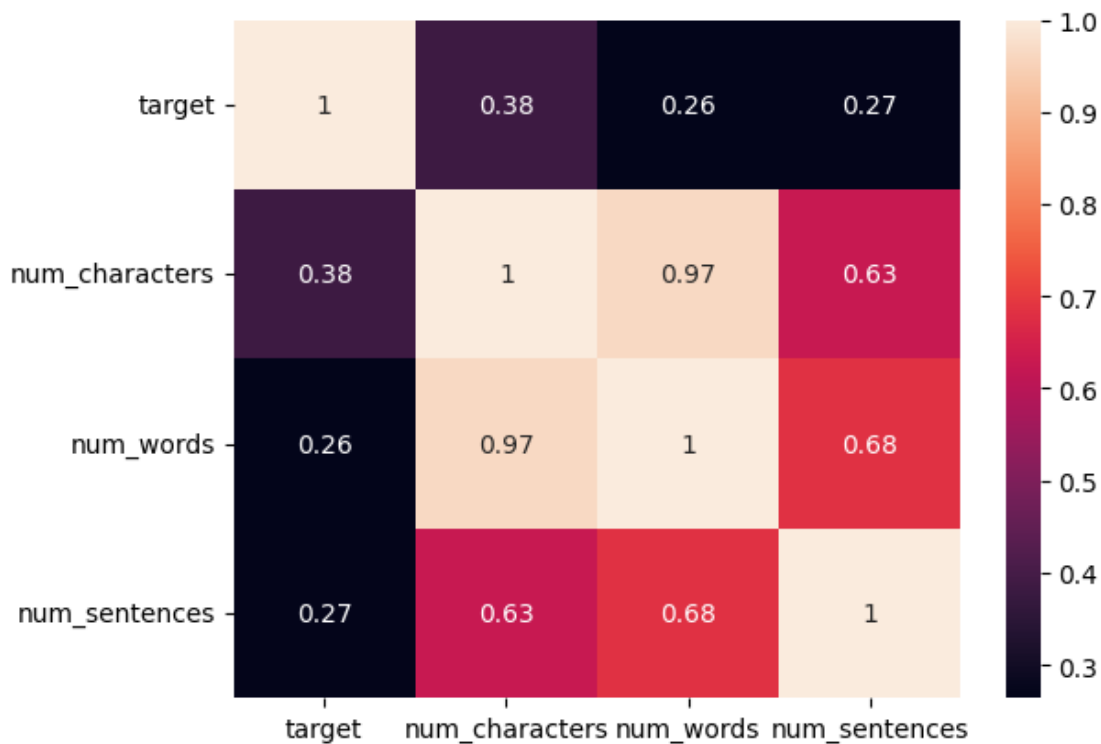
```
[185]: from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
data['target']=encoder.fit_transform(data['target'])
```

```
[186]: data['target'].value_counts()
```

```
[186]: target
0    4516
1     653
Name: count, dtype: int64
```

```
[187]: sns.heatmap(data.corr(numeric_only=True),annot=True)
```

```
[187]: <Axes: >
```



Preprocessing

```
[188]: nltk.download('stopwords')
```

[nltk_data] Downloading package stopwords to /usr/share/nltk_data...

[nltk_data] Package stopwords is already up-to-date!

```
[188]: True
```

```
[189]: from nltk.corpus import stopwords
import string
from nltk.stem.porter import PorterStemmer
ps = PorterStemmer()
```

```
[190]: def transform_text(text):
    # Convert text to lowercase
    text = text.lower()
    # Tokenize the text
    tokens = nltk.word_tokenize(text)

    # Remove punctuation and stopwords, and perform stemming
    transformed_tokens = [ps.stem(token) for token in tokens if token.isalnum()
    ↪ and token not in stopwords.words('english')]
```

```
return " ".join(transformed_tokens)
```

```
[191]: data['Email_transformed']=data['Email'].apply(transform_text)
```

```
[192]: data.head()
```

```
[192]:
```

	target	Email	num_characters	\
0	0	Go until jurong point, crazy.. Available only ...	111	
1	0	Ok lar... Joking wif u oni...	29	
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	
3	0	U dun say so early hor... U c already then say...	49	
4	0	Nah I don't think he goes to usf, he lives aro...	61	

	num_words	num_sentences	Email_transformed
0	23	2	go jurong point avail bugi n great world la e ...
1	8	2	ok lar joke wif u oni
2	37	2	free entri 2 wkli comp win fa cup final tkt 21...
3	13	1	u dun say earli hor u c already say
4	15	1	nah think goe usf live around though

Feature Selection

```
[212]: from sklearn.feature_extraction.text import CountVectorizer,TfidfVectorizer
tfidf = TfidfVectorizer()
X = tfidf.fit_transform(data['Email_transformed']).toarray()
```

```
[213]: y=data['target']
```

```
[214]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix
import time
```

```
[215]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=101)
```

```
[216]: model=MultinomialNB()
```

```
[217]: start_time = time.time()
model.fit(X_train, y_train)
train_time = time.time() - start_time
print("Training Time:", train_time)
```

Training Time: 0.05539870262145996


Model Evaluation

```
[218]: y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9451966473243069

```
[219]: conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix,annot=True)
```

[219]: <Axes: >



```
[220]: from sklearn.naive_bayes import GaussianNB, BernoulliNB
```

```
[221]: nb_classifier = GaussianNB()

nb_classifier.fit(X_train, y_train)
```

[221]: GaussianNB()

```
[222]: y_pred = nb_classifier.predict(X_test)


# Calculate accuracy
```

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.8594455190199871

```
[223]: # Generate confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix,annot=True)
```

[223]: <Axes: >



```
[224]: nb_classifier = BernoulliNB()

# Train the classifier on the training data
nb_classifier.fit(X_train, y_train)

# Make predictions on the testing data
y_pred = nb_classifier.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```


Accuracy: 0.9593810444874274

Conclusion

Best Model is MultinomialNB with accuracy of 97.43 %

Prediction on test data

```
[225]: actual_inverse = encoder.inverse_transform(y_test)
       predicted_inverse = encoder.inverse_transform(y_pred)

       # Creating a new DataFrame with inverse transformed values
       actual_vs_predicted = pd.DataFrame({'Actual': actual_inverse, 'Predicted':
       ↪ predicted_inverse})
```

```
[226]: actual_vs_predicted.head()
```

```
[226]:   Actual Predicted
0      ham         ham
1      ham         ham
2     spam         spam
3      ham         ham
4      ham         ham
```

Prediction on user input

```
[227]: import joblib
```

```
[238]: model_1=joblib.load('/kaggle/working/nb_classifier.joblib')
       tfid=joblib.load('/kaggle/working/tfid.joblib')
```

```
[ ]:
```

```
[239]: input_email = input('Enter Email: ')

       # Preprocess the user input
       preprocessed_input = [transform_text(input_email)]

       # Vectorize the preprocessed input using the already fitted vectorizer
       preprocessed_input_vectorized = tfid.transform(preprocessed_input)

       # Make predictions
       prediction = model_1.predict(preprocessed_input_vectorized)

       # Print the result
       if prediction == 0:
           print('ham')
       else:
           print('spam')
```

```
Enter Email: input_email =input('Enter Email') # Preprocess the user input
preprocessed_input = [transform_text(input_email)] # Vectorize tDon't miss out
on our exclusive offer! Get 50% off your next purchase when you shop with us
today. Limited time onlyhe preprocessed input preprocessed_input_vectorized =
vectorizer.fit_transform(preprocessed_input) print('') # Make predictions
prediction = model_1.predict(preprocessed_input_vectorized) if prediction==0:
print('ham') else:      print('spam')
```

ham

saving best model

```
[233]: from joblib import dump

# Assuming your model is named 'model'
dump(nb_classifier, 'nb_classifier.joblib')
dump(tfidf, 'tfidf.joblib')
```

```
[233]: ['tfidf.joblib']
```

```
[ ]: dump(transform_text, 'preprocess_fn.joblib')
```

```
[ ]:
```